

StoryEval: Adaptive Intra- and Inter-Segment Evaluation for Story-Grounded Video Generation

Yajing Xu
Zhejiang University
Hangzhou, Zhejiang, China
yajingxu@zju.edu.cn

Yiwei Wang
East China Normal University
Shanghai, China
ywwangworkhard@gmail.com

Haonan He
Nanjing University of Posts and
Telecommunications
Nanjing, Jiangsu, China
haonanhe01@gmail.com

Yichi Zhang
Zhejiang University
Hangzhou, Zhejiang, China
zhangyichi.each@zju.edu.cn

Ren He
University of Emergency
Management
Langfang, Hebei, China
heren@st.cidp.edu.cn

Xin Cao
Ocean University of China
Qingdao, Shandong, China
caoxin@stu.ouc.edu.cn

Huan Gao[✉]
sapiens company
Singapore
huan.gao@sapiens-ai.io

Wen Zhang[✉]
Zhejiang University
Hangzhou, Zhejiang, China
zhang.wen@zju.edu.cn

Abstract

Recent advances in video generation are making applications such as short-drama production and literary adaptation increasingly feasible, but current generators still produce clips much shorter than most target narratives. Long stories are therefore generated segment by segment, creating two evaluation needs: whether each segment realizes source-specified events and whether adjacent segments maintain visual and narrative continuity. Existing benchmarks largely score isolated clips and overlook this joint requirement. To bridge this gap, we introduce StoryEval, a benchmark for adaptive intra- and inter-segment evaluation of story-grounded video generation. StoryEval derives shared event trajectories and boundary-aware segment plans from source stories, combines Core metrics for cross-setting comparison with Adaptive metrics for story- and genre-specific risks, and uses a multimodal story asset graph for consistent generation and reference-guided vision-language model (VLM) evaluation. We instantiate StoryEval on 60 short-drama excerpts spanning 12 genres and eight Chinese and world classics. Evaluation of six generators shows that scene and style realization are relatively strong, whereas key-event realization, narrative transitions, and long-arc continuity remain challenging; open-weight systems also exhibit declining inter-segment scores along sequences of up to 10 segments. Data and code are available at https://gitee.com/kawi_t2v_benchmark/story-eval.

CCS Concepts

• **Computing methodologies** → **Computer vision**; *Artificial intelligence*; • **Information systems** → *Multimedia information systems*.

Keywords

video generation, story-grounded generation, multimodal benchmark, multi-segment continuity, video evaluation

ACM Reference Format:

Yajing Xu, Yiwei Wang, Haonan He, Yichi Zhang, Ren He, Xin Cao, Huan Gao, and Wen Zhang. 2027. StoryEval: Adaptive Intra- and Inter-Segment Evaluation for Story-Grounded Video Generation. In *Proceedings of ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD '27)*. ACM, New York, NY, USA, 24 pages.

1 Introduction

Recent advances in text-to-video (T2V) generation have substantially improved visual quality, motion realism, and prompt alignment [10, 25]. These advances are bringing applications such as short-drama production, literary adaptation, and educational storytelling closer to practical use. However, the duration produced by current generators remains much shorter than that required by most narrative videos. In practice, a long story is therefore planned and generated segment by segment before being assembled into a complete sequence.

This workflow changes what generation quality means. A successful system must first realize the intended content within each segment, including the required characters, actions, and events. It must also preserve or correctly update visual and narrative state across segment boundaries. A clip may look realistic in isolation but still fail as part of a story if it omits a key event, changes a character unexpectedly, or breaks the connection to the next segment.

These requirements are especially important in applications grounded in a source story, such as adapting literary works to video or producing short dramas from scripts. In these settings, visual plausibility alone is not enough. The generated sequence should realize the core events specified by the source and maintain their intended progression. Evaluation must therefore address two complementary questions: whether each segment faithfully

Table 1: Comparison of StoryEval with representative video generation benchmarks. Evaluation units are grouped by temporal structure: isolated clip, ordered multi-clip, hierarchical movie units, or event-planned multi-segment evaluation. A dash means cross-clip evaluation is not a primary target.

Benchmark	Evaluation unit	Source input	Intra-clip focus	Inter-clip focus	Evaluation method
VBench / VBench++ [8, 9]	Isolated clip	Text prompt	General quality, motion, and alignment	–	Fixed automatic metrics
EvalCrafter [15]	Isolated clip	Text prompt	Visual, content, motion, and alignment quality	–	Fixed automatic metrics and human ratings
FETV [17]	Isolated clip	Fine-grained prompt	Content, attribute, and temporal adherence	–	Fixed automatic metrics
T2V-CompBench [23]	Isolated clip	Compositional prompt	Binding, relation, action, motion, and counting	–	Fixed automatic metrics
StoryBench [3]	Ordered multi-clip	Ordered action captions	Action realization and visual quality	Action continuity	Fixed automatic metrics
MovieBench [26]	Hierarchical movie units	Synopsis and hierarchical descriptions	Shot- and scene-level quality	Movie-level coherence	Automatic metrics with character bank
StoryEval	Event-planned multi-segment	Source story and event plan	Event and entity realization	Visual and narrative continuity	Reference-guided VLM with Core and Adaptive metrics

presents its assigned story content, and whether adjacent segments form a coherent narrative sequence.

Existing benchmarks largely evaluate isolated clips through visual quality, prompt alignment, composition, physical plausibility, temporal dynamics, safety, or human preference [2, 7–9, 15, 17, 18, 23, 28]. StoryBench and MovieBench move toward continuous actions and movie-level structure [3, 26]. As summarized in Table 1, however, they do not jointly evaluate source-grounded event realization within segments and boundary-aware continuity across segments. Consequently, existing benchmark rankings may reward visually appealing clips while overlooking whether they collectively tell the intended story.

To bridge this gap, we introduce StoryEval, a benchmark for adaptive intra- and inter-segment evaluation of story-grounded video generation. StoryEval converts each source excerpt into a shared event trajectory and boundary-aware segment plans. These plans provide common generation and evaluation units across systems. Core metrics assess recurring requirements and support comparison across settings, while Adaptive metrics capture risks specific to a story or genre. StoryEval further constructs a source-grounded multimodal story asset graph that links event nodes to character, object, and scene cards, supporting consistent generation and reference-guided vision-language model (VLM) evaluation [7, 19].

We instantiate StoryEval on 60 short-drama excerpts from eight Chinese and world classics, covering 12 genres. This setting provides event-dense narratives with recurring characters, objects, and locations, while the protocol can also be applied to other novels, scripts, and narrative texts. Evaluation of six practical systems reveals a consistent pattern. Models perform relatively well on scene and style realization, but struggle with key events, narrative transitions, and long-arc continuity. Continuous-frame boundaries are easier than narrative transitions, and the inter-segment scores of open-weight systems decline along sequences of up to 10 segments even when their individual clips remain visually plausible.

In summary, our contributions are:

- **Resource.** We introduce StoryEval and instantiate it on 60 story-grounded multi-segment instances spanning 12 genres and eight Chinese and world classics. Data and code are

available at https://gitee.com/kawi_t2v_benchmark/story-eval.

- **Protocol.** We develop a unified protocol for event-to-segment planning, Core+Adaptive metric construction, and reference-guided VLM evaluation. It supports both intra-segment event diagnosis and boundary-aware inter-segment evaluation.
- **Diagnosis.** Evaluation of six open-weight and closed models reveals that strong clip-level appearance does not imply successful storytelling. Key-event realization, narrative transitions, and long-arc continuity remain major bottlenecks, with human studies supporting both metric selection and VLM-based scoring.

2 Related Work

Video generation for narrative content. Diffusion- and transformer-based systems increasingly support realistic motion, image conditioning, continuation, and multi-shot generation [10, 25]. Practical applications can therefore combine T2V, text-to-image (T2I), and image-to-video (I2V) components to generate a narrative sequence segment by segment. However, systems differ in their native clip duration, conditioning interface, and continuation mechanism. Comparing them as narrative generators requires a common source-grounded plan rather than independently authored prompts.

Clip-centric video evaluation. VBench, VBench++, EvalCrafter, and FETV provide broad assessment of visual quality, motion, temporal consistency, and prompt alignment [8, 9, 15, 17]. Other benchmarks isolate composition, physical plausibility, motion dynamics, safety, or human preference, while VideoScore learns fine-grained automatic feedback for generated videos [2, 7, 14, 18, 23, 28]. Despite this increasing granularity, the dominant evaluation unit remains a prompt–clip pair. It does not assign source events to individual segments or specify which visual and narrative states should persist across their boundaries.

Story and long-video evaluation. StoryBench moves beyond isolated clips by evaluating continuous visualization from ordered




Figure 1: Overview of StoryEval. From a source excerpt, the benchmark builds a shared event trajectory and boundary-aware segment plans, runs model-native sequential generation, and scores Intra/Inter items with Core+Adaptive metrics via a multimodal story asset graph and reference-guided VLM judging.

action captions, while MovieBench provides hierarchical annotations at the movie, scene, shot, and character levels [3, 26]. Recent multi-shot benchmarks such as MSVBench further assess cross-shot coherence with hierarchical scripts and reference images [22]. These efforts advance long-form evaluation, but they do not jointly represent which source-story events each segment should realize and what state should be preserved or updated at each boundary. StoryEval addresses this gap with source-grounded event plans, boundary-typed continuity evaluation, and Core+Adaptive metrics supported by a multimodal story asset graph.

3 StoryEval Framework

3.1 Task and Framework Overview

StoryEval evaluates a complete story-grounded generation system under a common story plan. Given a source excerpt x and its genre g , the benchmark first constructs an ordered event trajectory $N = (n_1, \dots, n_T)$ and boundary-aware segment plans $V = (v_1, \dots, v_M)$. Every system receives the same event ownership and boundary requirements, although the plans are adapted to its native generation interface. Segments are evaluated for event and entity realization; adjacent-segment boundaries are evaluated for the visual or narrative state they should preserve or update.

As shown in Figure 1, StoryEval separates benchmark preparation, system generation, and evaluation. Before testing any system, it fixes the plans V , observable scoring items Q , and a multimodal story asset graph G_s , exposed as two views: G_s^{gen} for prompting cues and G_s^{eval} for reference-guided VLM judging. For a system F_θ , the canonical plan v_j is converted into its model-specific input p_j^θ with cues from G_s^{gen} , producing the sequence $Y = (y_1, \dots, y_M)$. A VLM judge then returns

$$R = \phi_{\text{judge}}(Y, Q, G_s^{\text{eval}}),$$

including normalized segment and boundary scores, item-level evidence, and failure diagnoses. This design compares complete generation systems under the same source-grounded plan rather than testing their independent story-planning ability.

3.2 Narrative-to-Segment Planning and Generation

Event trajectory. Directly prompting with a long passage obscures event ownership, whereas length-based splitting can break causal and emotional progression. StoryEval prompts a large language model (LLM) to extract event nodes

$$n_i = (\text{id}_i, \text{name}_i, \text{text}_i),$$

where text_i is a verbatim source span and nodes preserve the original narrative order. Nodes cover visible actions as well as decisions, revelations, relationship or emotional changes, scene transitions, and outcomes that alter the narrative state.

Segment plans. Planning proceeds in two LLM stages: first determine the ordered segment-to-node mapping, then instantiate generation-ready plans. Every node must be covered without re-ordering; one segment may cover consecutive nodes, and a complex node may span consecutive segments only through distinct, non-overlapping sub-event ranges. Each plan contains source coverage, a filmable scene description, character/scene/action/emotion/camera fields, and boundary metadata. We target 5–10 seconds per segment, matching the common native duration range of current video generation systems [10, 25] while keeping one visually observable micro-event per clip. Optional cinematic details guide generation but are not treated as mandatory source-faithfulness evidence. Appendix D provides the full schema, prompt constraints, and generation protocol.

Boundary-aware generation. To avoid treating intended scene or time changes as continuity failures, StoryEval distinguishes two relations between adjacent segments. A *continuous-frame anchor* links segments that extend the same visual or action state, whereas a *narrative transition* links segments across an intended change in time, location, viewpoint, or narrative line. Accordingly, segment j is generated as

$$y_j = \begin{cases} F_{\theta}^{\text{T2V}}(p_j^{\theta}), & j = 1 \text{ or } b_j = \text{transition}, \\ F_{\theta}^{\text{I2V}}(p_j^{\theta}, f_{j-1}^{\text{end}}), & j > 1 \text{ and } b_j = \text{continuous}. \end{cases}$$

Here, f_{j-1}^{end} is the actual last frame of y_{j-1} . For an I2V-only system, a model-family T2I component generates the initial frame for the first segment and for segments following narrative transitions. Generation-side character descriptions are retrieved from G_s^{gen} . Model-specific adaptation and inference settings are reported with the experiments.

3.3 Core and Adaptive Metric Design

Evaluating story-grounded video requires balancing cross-setting comparability with genre-specific diagnostic coverage. A fully fixed metric set can miss failure modes emphasized by particular genres, whereas manually defining metrics for individual segments weakens reusability and cross-model comparison. We therefore construct a shared candidate metric library as a four-level ontology of *domain*, *object*, *attribute*, and *metric*. The library covers 8 segment-level and 9 boundary-level evaluation-object categories, including event and key-object realization within segments and character, action, and causal continuity across segments.

Core metrics. Drawing on prior video-generation benchmarks [8, 9, 15, 17], we manually select Core metrics according to three criteria: coverage of fundamental generation quality, visual observability, and cross-genre comparability. Segment-level Core evaluation uses one metric for each of the eight evaluation objects: Character Identity Accuracy (CIA), Key Event Accuracy (KEA), Scene & Element Consistency (SEC), Key Object Presence (KOP), Emotion Match Rate (EMR), Subject Visibility (SV), Costume & Environment Consistency (CEC), and Visual Stability (VS). Boundary-level

Core evaluation retains Transition Smoothness (TS), Visual Texture Continuity (VTC), and Splice Quality Stability (SQS). Selection rationales are listed in Appendix Table 11. The resulting Core set is fixed across genres, stories, and systems.

Adaptive metrics. Adaptive metrics complement the fixed Core set by targeting requirements emphasized within each genre. We select them before video generation so that the evaluation criteria cannot adapt to observed model outputs. For each genre, an LLM extracts source-grounded evidence units from event nodes and planned boundaries, maps them to ontology paths, and assigns $L(e) \in \{0, 1, 2, 3\}$ by importance and visual observability. Let $C_g(m)$ be the number of stories in genre g with valid evidence ($L \geq 2$) for metric m , and let $\bar{D}_g(m)$ be its mean valid unit-level support. We discard metrics with insufficient coverage ($C_g(m) < K_{\text{cov}}$, $K_{\text{cov}} = 3$) and keep the $N_d = 5$ remaining metrics with the largest $\bar{D}_g(m)$:

$$M_{g,d}^{\text{adaptive}} = \text{Top}_{N_d}(\{m \in M_d : C_g(m) \geq K_{\text{cov}}\}; \bar{D}_g),$$

where $\text{Top}_{N_d}(\cdot; \bar{D}_g)$ returns the N_d highest-scoring metrics under $\bar{D}_g$, and ties are broken by a fixed ontology order. The selected metrics are then combined with the fixed Core set for every story in genre g . Human validation of this protocol is in Section 5.5; ontology details and the ranking algorithm are in Appendix E.

Evaluation item instantiation. Before generation, a constrained LLM converts the fixed Core and Adaptive metrics into unit-specific evaluation items Q , grounded in the source excerpt, covered event range, and plan fields. The metric inventory is frozen: the LLM may neither add, delete, nor rename metrics. Each item specifies an evaluation goal, visually observable checks, and a 0/1/2 rubric; unobservable mental states and optional cinematic elaborations are excluded.

3.4 Story Asset Grounding and VLM Scoring

StoryEval relies on a VLM judge to score the instantiated evaluation items, following recent automatic evaluation practice for generated media [7, 12, 16, 19]. In multi-character, multi-object, and narratively structured segments, however, a VLM may attribute an action to the wrong character or misrecognize the category, form, or period style of a key object. Such recognition errors can be conflated with genuine generation failures and therefore distort the resulting scores.

To mitigate this risk, StoryEval constructs a story-level multi-modal reference graph $G_s = (V_s, E_s)$ for each source story. The graph stores event nodes together with character, object, and scene cards, and is indexed by segments through the event nodes they cover. As introduced in Section 3.1, G_s^{gen} exposes only character identity and node-specific visible-state descriptions for prompting, while G_s^{eval} retrieves segment- or boundary-relevant cards as recognition aids for VLM judging. Neither view is a pixel-level answer key or a full T2V conditioning signal; evaluation items, rubrics, and scores remain withheld from the generator.

Character exemplars and object references serve only as recognition aids (Appendix F).

Scoring protocol. An evaluation unit u is either a segment v_j or an adjacent-segment boundary (v_j, v_{j+1}) . For metric m , the VLM scores every item $q \in Q(u, m)$ on $\{0, 1, 2\}$. Dividing by the maximum

possible item sum yields

$$\hat{S}(u, m) = \frac{1}{2K_{u,m}} \sum_{q \in Q(u, m)} s(q), \quad K_{u,m} = |Q(u, m)|. \quad (1)$$

Story scores average over applicable units:

$$S_c(m) = \frac{1}{|U_c(m)|} \sum_{u \in U_c(m)} \hat{S}(u, m). \quad (2)$$

System-level results give all 60 stories equal weight. Boundary scoring uses the last two seconds of y_j and the first two seconds of y_{j+1} , with continuous-frame anchors (CFA) checking visual and state continuity and narrative transitions (NT) checking causal and emotional inheritance without penalizing intended changes.

4 StoryEval Instantiation

The protocol in Section 3 is general. To make it a reusable benchmark, we instantiate it on classical narrative excerpts that provide stable source grounding, event-dense plots, and recurring entities. We draw from the Four Great Classical Novels of Chinese literature and four world classics, so that evaluation does not depend on a single narrative tradition. Contemporary short-drama genre labels are used only as diagnostic dimensions; they do not redefine the underlying source texts.

Genre coverage and excerpt selection. We organize the testbed around 12 short-drama genres: Drama (Dram), Stratagem (Strat), Action (Act), Fantasy (Fant), Adventure (Adv), Comedy (Com), Suspense (Susp), Historical Romance (H-Rom), Female Growth (FGrow), Palace Intrigue (Pal-Int), Patriotism (Patr), and Supernatural Tales (Sp-Tales). Operational definitions are in Appendix H. Each genre contains five excerpts jointly selected from Chinese and world classics, yielding 60 excerpts with clear narrative function, visually salient events, and compact event closure.

Scale, release, and quality control. Figure 2 summarizes the resulting resource. Across the 60 stories, StoryEval contains 1,347 event nodes and 1,488 planned segments (mean 22.4 nodes and 24.8 segments per story). Among the 1,428 adjacent-segment boundaries in these plans, 1,052 are CFA and 376 are NT. The complete plans define the released benchmark units. For comparative diagnosis under a shared story arc, the experiments in Section 5 evaluate the most central contiguous block of up to 10 planned segments in each story. This diagnostic subset contains 584 segments and 524 boundaries (394 CFA and 130 NT). Source excerpts, event and segment plans, evaluation items, story asset graphs, and scoring code are available at https://gitee.com/kawi_t2v_benchmark/story-eval.

Construction follows stage-specific criteria in Appendix H; flagged outputs are revised or regenerated, but we do not exhaustively audit every boundary or evaluation item. During adaptive-metric validation (Section 5.5), annotators consult source text, nodes, and segment and boundary plans and report no systematic inconsistencies in the sampled units; that section also reports VLM–human agreement.




Figure 2: StoryEval instantiation scale: 12 genres $\times$ 60 excerpts yield 1,488 planned segments from eight Chinese and world classics.

5 Experiments and Analysis

5.1 Setup and Systems

Systems and scoring. We evaluate six systems under a shared StoryEval plan: open-weight JoyAI-Echo [6] (JoyAI), LTX-2.3 [13] (LTX), and HunyuanVideo-1.5 [24] (Hunyuan); and closed models Seedance 2.0 [4] (Seedance), HappyHorse 1.1 [1] (HappyHorse), and Grok Imagine Video 1.5 [27] (Grok); adaptations are in Appendix A. Unless a version tag is needed for disambiguation, we use these short names in the text and figures below. As in Section 4, each story uses its central contiguous block of up to 10 segments (584 clips; 524 boundaries). All systems generate at 720p; I2V models reuse the previous end frame only under CFA, while JoyAI is T2V-only. Qwen3.6-27B [20] scores all items under the reference-guided protocol; we report story-level macro-averages over the 60 stories.

5.2 Main Results

Table 2 summarizes system-level performance under the full Core+Adaptive protocol. HappyHorse leads Overall (Intra 0.92, Inter 0.89), with Grok close behind; Seedance remains strong on Intra (0.90) but is matched on Inter by open-weight JoyAI (0.86). Among open-weight systems, JoyAI is the strongest overall, while LTX and Hunyuan lag on both Intra and Inter.

Surface appearance is largely solved; storytelling is not. Scene scores stay high across systems (often ≥ 0.90), and Style-I is similarly strong, while Style-B drops more (0.75–0.88). Character and object fidelity, key-event realization, and action continuity leave substantial headroom: Event-I spans only 0.71–0.90, and Action-B is the weakest shared boundary column (0.64–0.82). For closed models and Hunyuan, Intra Overall also exceeds Inter, with the sharpest continuity drop on stability and action (e.g., Hunyuan Stability 0.70 $\rightarrow$ 0.56). A system can thus look cinematic on individual clips while still failing to carry plot state across the story arc. This is the gap that clip-centric benchmarks tend to miss.

Table 2: Main StoryEval results under the full Core+Adaptive protocol (story-level means in [0, 1]). I/B: intra-video / boundary-level inter-video. Bold: best in each column. System names follow Section 5.1.

Model	Character		Object		Scene		Emotion		Style		Stability		Event	Action	Causal	Overall	
	I	B	I	B	I	B	I	B	I	B	I	B				I	B
Hunyuan	0.74	0.76	0.72	0.74	0.93	0.92	0.77	0.84	0.89	0.75	0.70	0.56	0.71	0.64	0.83	0.78	0.75
LTX	0.76	0.80	0.73	0.76	0.93	0.94	0.81	0.83	0.89	0.85	0.92	0.88	0.76	0.74	0.86	0.83	0.84
JoyAI	0.82	0.79	0.76	0.81	0.93	0.93	0.84	0.91	0.93	0.82	0.92	0.96	0.79	0.74	0.90	0.86	0.86
Grok	0.88	0.84	0.86	0.85	0.96	0.93	0.91	0.90	0.97	0.86	0.92	0.95	0.88	0.77	0.89	0.91	0.87
Seedance	0.87	0.84	0.83	0.83	0.96	0.92	0.90	0.87	0.96	0.83	0.94	0.91	0.86	0.73	0.90	0.90	0.85
HappyHorse	0.91	0.88	0.87	0.88	0.97	0.94	0.92	0.93	0.97	0.88	0.93	0.94	0.90	0.82	0.90	0.92	0.89

5.3 Metric Diagnosis

To locate where StoryEval is hard, we decompose Table 2 into shared Core metrics and evidence-selected Adaptive metrics (Figure 3; per-system panels in Appendices B and C).

Core overview. Figure 3(a) shows that the closed>open ranking is stable across genres. Genre mainly changes *how severe* the failures are, not *who* leads. Averaged over systems, Action, Fantasy, and Supernatural Tales are the hardest settings, while Suspense, Strata-gem, and Historical Romance remain comparatively easier. More importantly, **metric identity outweighs genre identity**. SEC, CEC, and SV are high for every system, whereas TS and KEA (with relatively weak KOP) remain bottlenecks even for closed models. SQS mainly isolates Hunyuan rather than revealing a shared failure mode. LTX is additionally weaker on CIA, EMR, and SV. **Core difficulty therefore concentrates on key-event realization and transition smoothness**, with SQS as an open-weight-specific limitation.

Genre-adaptive diagnosis. Adaptive metrics sharpen the same picture. Figure 3(b) preserves the system ranking but enlarges closed-open gaps on Fantasy, Sp-Tales, and Action (about +0.10–0.11). In Figure 3(c), Action Accuracy (AAC), Action Flow (AFL), and Object Interaction (OIN) repeatedly form the lowest Adaptive stack. Spatial Continuity (SCT), Atmosphere Alignment (AAL), and Visual Consistency (VCO) remain comparatively strong when selected. Thus Adaptive scoring does not merely restate Core scene and style strength. **It exposes failures in executing actions, inheriting action state, and interacting with objects.** HappyHorse stays relatively robust on these hard metrics, while Hunyuan keeps the same ordering at a lower absolute level. Appendix C shows that Grok and Seedance resemble HappyHorse at a slightly lower tier, with LTX and JoyAI in between.

5.4 Long-arc Continuity

Overall Inter scores alone leave several continuity factors entangled. We therefore analyze boundary type, segment index, and multi-hop character reuse separately.

Boundary type (CFA vs NT). Table 3 splits inter-video scores by CFA and NT. **Continuous-frame anchors are consistently easier than narrative transitions.** CFA Overall is at least as high as NT for every system, with the largest gaps on Grok (0.90 → 0.82) and Seedance (0.87 → 0.81). HappyHorse leads CFA (0.91); on NT,

Table 3: Inter-video continuity by boundary type (CFA vs. NT; story-level means in [0, 1]). Bold: best within each type. System names follow Section 5.1.

Model	Type	Char	Obj	Scn	Emo	Sty	Stab	Act	Caus	All
Hunyuan	CFA	0.76	0.73	0.93	0.85	0.77	0.57	0.67	0.83	0.76
	NT	0.78	0.77	0.86	0.87	0.70	0.52	0.54	0.85	0.73
LTX	CFA	0.80	0.75	0.97	0.84	0.88	0.89	0.77	0.86	0.85
	NT	0.82	0.79	0.86	0.80	0.76	0.87	0.66	0.90	0.81
JoyAI	CFA	0.80	0.79	0.95	0.90	0.81	0.96	0.76	0.90	0.86
	NT	0.80	0.87	0.87	0.93	0.86	0.95	0.74	0.90	0.86
Grok	CFA	0.87	0.86	0.95	0.90	0.89	0.95	0.81	0.90	0.90
	NT	0.77	0.83	0.87	0.90	0.75	0.92	0.70	0.87	0.82
Seedance	CFA	0.85	0.83	0.94	0.87	0.87	0.94	0.76	0.91	0.87
	NT	0.84	0.79	0.86	0.88	0.72	0.83	0.66	0.88	0.81
HappyHorse	CFA	0.89	0.87	0.96	0.94	0.91	0.94	0.85	0.90	0.91
	NT	0.87	0.83	0.87	0.90	0.81	0.94	0.77	0.90	0.86

HappyHorse and JoyAI tie (0.86). Action remains the weakest shared column under both types ($CFA \leq 0.85$; $NT \leq 0.77$). This asymmetry is expected under StoryEval’s design. CFA mainly asks for visual and state inheritance. NT requires causal and emotional carry-over without punishing intended scene or time changes, and is therefore a harder storytelling demand.

Segment-order decay and character drift. Figures 4 and 5 ask whether competence persists over a sequence of up to 10 segments. Intra-Closed is nearly flat (0.95 → 0.93), Intra-Open declines mildly (0.87 → 0.83), and **Inter-Open decays most clearly** (0.82 → 0.77). The pattern suggests that open-weight systems can still look strong on late Intra clips while gradually losing cross-segment state. Character appearance tells the same story under CFA reuse. CLIP [21] similarity versus reuse lag (HappyHorse vs. LTX) falls from about 0.84 at lag 1 to 0.76–0.77 by lag 4. HappyHorse means at reuse lags 2, 3, and 5 (0.844, 0.801, and 0.764) further show that **even the strongest closed system drifts under multi-hop identity reuse**.

5.5 Protocol Validity

We next validate Adaptive metric selection and reference-guided VLM scoring against human judgments.

Human validation of adaptive metrics. For each of 12 genres, annotators score two Intra and two Inter Adaptive metrics (three evidence units for Q1; three items for Q2), yielding 144 tasks ×




Figure 3: Metric diagnosis on StoryEval (story-level means in $[0, 1]$). (a) Core heatmaps by genre and by metric. (b) Adaptive Intra and Inter genre overviews. (c) Adaptive metric stacks for HappyHorse (top) and Hunyuan (bottom); absent colors mean the metric is not selected. Abbreviations follow Sections 5.1, 4, and 3.3.




Figure 4: Mean Core score versus segment index, aggregated by access type. Solid: Intra; dashed: Inter.

three annotators on $\{0, 1, 2\}$. Q1 asks whether the evidence chain justifies selecting the metric; Q2 asks whether the item is scorable from the video window. Table 4 summarizes the study. Mean scores are 1.90 for Q1 and 1.94 for Q2. Majority =2 holds on 93.1% and 95.1% of tasks, and Exact-3 on 86.8% and 90.3%. Ordinal Krippendorff’s α [11] is only moderate (0.51 for Q1 and 0.45 for Q2) under heavy score-2 imbalance. We therefore emphasize agreement rates rather than α alone. Annotators who consult source text and plans also report no systematic inconsistencies in the sampled units (Section 4). **Adaptive metrics are thus human-supported before generation.** They are not post-hoc labels fitted to model outputs.

VLM–human agreement. On 432 items (1,548 checks) spanning 12 genres, 6 systems, and both Intra and Inter domains, three annotators use the same rubric as Qwen3.6-27B [20]. Gold is the score agreed by at least two annotators. Human–human pairwise exact agreement is 58.9% (MAE 0.47; ordinal $\alpha=0.42$ [11]). This underscores that story-grounded scoring is inherently subjective even for humans. Table 5 compares + ref and – ref against this gold label.




Figure 5: CFA appearance drift: CLIP [21] similarity versus reuse lag for HappyHorse and LTX (n at each lag).

Table 4: Human validation of adaptive-metric selection (3 annotators; 144 tasks on $\{0, 1, 2\}$). Q1: evidence-chain support; Q2: item judgeability. Mean: annotator-then-task average; $\% (=2)$: majority=2; Exact-3: full agreement; α : ordinal Krippendorff. [†]Low Intra-Q2 α despite high Exact-3 reflects score-2 dominance.

Split	Q	Mean	$\% (=2)$	Exact-3	α
All	Q1	1.90	93.1	86.8	0.51
	Q2	1.94	95.1	90.3	0.45
Intra	Q1	1.90	94.4	87.5	0.53
	Q2	1.96	97.2	90.3	0.19 [†]
Inter	Q1	1.90	91.7	86.1	0.50
	Q2	1.92	93.1	90.3	0.58




Figure 6: Case study of long-arc and reference-sensitive failures. (a) CFA appearance drift of Lin Daiyu from reference V_5 under pose and framing changes ($0.91 \rightarrow 0.77$). (b) Intra item: without references the VLM confuses Zhang Fei with Guan Yu; with references the score drops from 2 to 0. (c) Inter item on Odysseus railing (V_{28} – V_{29}): – ref invents continuity, while + ref correctly rejects it (likewise $2 \rightarrow 0$).

Table 5: Human–human consistency and VLM–human agreement on $\{0, 1, 2\}$ checks (gold: ≥ 2 -annotator agreement; 1,453/1,548). + ref/– ref: with/without reference cards. Exact: % exact match; MAE. Last column: Human–Human α ; VLM rows: Spearman ρ (+ref/– ref).

	Exact	± 1	MAE	ρ or α
<i>Human–Human</i>				
Pairwise avg.	58.9	–	0.47	$\alpha=0.42$
Exact-3 / ≥ 2 agree	41.4 / 93.9	–	–	–
<i>VLM vs. Human gold (+ref/– ref)</i>				
All	65.6 / 64.3	93.1 / 90.5	0.41 / 0.45	0.37 / 0.31
Intra	69.1 / 69.4	95.0 / 93.9	0.36 / 0.37	0.42 / 0.34
Inter	61.4 / 58.3	90.9 / 86.4	0.48 / 0.55	0.31 / 0.27
Core	69.3 / 68.8	94.6 / 92.1	0.36 / 0.39	0.43 / 0.39
Adaptive	61.2 / 59.1	91.3 / 88.7	0.48 / 0.52	0.30 / 0.22

Under + ref, exact agreement reaches 65.6% and within-one agreement reaches 93.1% (MAE 0.41), comparable to pairwise human agreement. Removing references raises Inter MAE from 0.48 to 0.55. **Reference-guided VLM scoring is therefore both necessary and competitive with human agreement on StoryEval’s check-level protocol.**

5.6 Case Study

Figure 6 illustrates the main diagnostic findings on three concrete instances. Panel (a) shows multi-hop CFA drift: Lin Daiyu remains locally plausible, yet similarity to reference V_5 drops under pose and framing changes (Section 5.4). Panels (b) and (c) show reference-sensitive judging: without cards, the VLM misattributes Zhang Fei’s table-press action and invents railing continuity across V_{28} – V_{29} ; with cards, both scores fall from 2 to 0. **Cinematic local clips can still lose long-arc identity, and reference cards are needed so that recognition errors are not counted as generation successes.**

6 Conclusion

In this paper, we introduced StoryEval, a benchmark for adaptive intra- and inter-segment evaluation of story-grounded video generation. Instantiated on classical short-drama excerpts spanning multiple genres, it couples shared event trajectories and boundary-aware segment plans with Core and Adaptive metrics, supported by a multimodal story asset graph for consistent generation and reference-guided scoring. Experiments on six open-weight and closed generators show that strong scene and style realization does not imply successful storytelling: key-event realization, narrative transitions, and long-arc continuity remain major bottlenecks. We release StoryEval to foster comparable progress on story-grounded video generation at https://gitee.com/kawi_t2v_benchmark/story-eval.

References

- [1] Alibaba. 2026. HappyHorse 1.1. Closed-source video generation API.
- [2] Hritik Bansal, Zongyu Lin, Tianyi Xie, Zeshun Zong, Michal Yarom, Yonatan Bitton, Chenfanfu Jiang, Yizhou Sun, Kai-Wei Chang, and Aditya Grover. 2025. VideoPhy: Evaluating Physical Commonsense for Video Generation. In *International Conference on Learning Representations*.
- [3] Emanuele Bugliarello, H. Hernan Moraldo, Ruben Villegas, Mohammad Babaeizadeh, Mohammad Taghi Saffar, Han Zhang, Dumitru Erhan, Vittorio Ferrari, Pieter-Jan Kindermans, and Paul Voigtlaender. 2023. StoryBench: A Multifaceted Benchmark for Continuous Story Visualization. In *Advances in Neural Information Processing Systems*.
- [4] ByteDance Seed. 2026. Seedance 2.0. https://seed.bytedance.com/en/seedance2_0.
- [5] DeepSeek-AI. 2026. DeepSeek-V4: Towards Highly Efficient Million-Token Context Intelligence. *arXiv preprint arXiv:2606.19348* (2026).
- [6] Echo Team @ Joy Future Academy, JD. 2026. *JoyAI-Echo: Pushing the Frontier of Long Audio-Visual Generation*. Technical Report. Joy Future Academy, JD. <https://github.com/jd-opensource/JoyAI-Echo>.
- [7] Xuan He, Dongfu Jiang, Ge Zhang, Max Ku, Achint Soni, Sherman Siu, Haonan Chen, Abhranil Chandra, Ziyang Jiang, Aaran Arulraj, Kai Wang, Quy Duc Do, Yuansheng Ni, Bohan Lyu, Yaswanth Narsupalli, Rongqi Fan, Zhiheng Lyu, Yuchen Lin, and Wenhui Chen. 2024. VideoScore: Building Automatic Metrics to Simulate Fine-grained Human Feedback for Video Generation. *CoRR* abs/2406.15252 (2024). doi:10.48550/arXiv.2406.15252
- [8] Ziqi Huang, Yinan He, Jiashuo Yu, Fan Zhang, Chenyang Si, Yuming Jiang, Yuanhan Zhang, Tianxing Wu, Qingyang Jin, Nattapol Chanpaisit, Yaohui Wang, Xinyuan Chen, Limin Wang, Dahua Lin, Yu Qiao, and Ziwei Liu. 2024. VBench: Comprehensive Benchmark Suite for Video Generative Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 21807–21818. doi:10.1109/CVPR52733.2024.02060
- [9] Ziqi Huang, Fan Zhang, Xiaojie Xu, Yinan He, Jiashuo Yu, Ziyue Dong, Qianli Ma, Nattapol Chanpaisit, Chenyang Si, Yuming Jiang, Yaohui Wang, Xinyuan Chen, Ying-Cong Chen, Limin Wang, Dahua Lin, Yu Qiao, and Ziwei Liu. 2026. VBench+: Comprehensive and Versatile Benchmark Suite for Video Generative Models. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 48, 3 (2026), 3268–3285. doi:10.1109/TPAMI.2025.3633890
- [10] Weijie Kong, Qi Tian, Zijian Zhang, Rox Min, Zuozhuo Dai, Jin Zhou, Jiangfeng Xiong, Xin Li, Bo Wu, Jianwei Zhang, Kathrina Wu, Qin Lin, Junkun Yuan, Yanxin Long, Aladdin Wang, Andong Wang, Changlin Li, Duojuan Huang, Fang Yang, Hao Tan, Hongmei Wang, Jacob Song, Jiawang Bai, Jianbing Wu, Jinbao Xue, Joey Wang, Kai Wang, Mengyang Liu, Pengyu Li, Shuai Li, Weiyan Wang, Wenqing Yu, Xincheng Deng, Yang Li, Yi Chen, Yutao Cui, Yuanbo Peng, Zhentao Yu, Zhiyu He, Zhiyong Xu, Zixiang Zhou, Zunnan Xu, Yangyu Tao, Qinglin Lu, Songtao Liu, Dax Zhou, Hongfa Wang, Yong Yang, Di Wang, Yuhong Liu, Jie Jiang, and Caesar Zhong. 2024. HunyuanVideo: A Systematic Framework for Large Video Generative Models. *CoRR* abs/2412.03603 (2024). doi:10.48550/arXiv.2412.03603
- [11] Klaus Krippendorff. 2018. *Content Analysis: An Introduction to Its Methodology* (4 ed.). Sage.
- [12] Max Ku, Dongfu Jiang, Cong Wei, Xiang Yue, and Wenhui Chen. 2024. VIEScore: Towards Explainable Metrics for Conditional Image Synthesis Evaluation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*. <https://aclanthology.org/2024.acl-long.663>
- [13] Lightricks. 2026. LTX-2.3. <https://huggingface.co/Lightricks/LTX-2.3>. See also Hareesh et al., arXiv:2601.03233.
- [14] Xinrang Ling, Chen Zhu, Meiqi Wu, Hangyu Li, Xiaokun Feng, Cundian Yang, Aiming Hao, Jiashu Zhu, Jiahong Wu, and Xiangxiang Chu. 2025. VM-Bench: A Benchmark for Perception-Aligned Video Motion Generation. *CoRR* abs/2503.10076 (2025). doi:10.48550/arXiv.2503.10076
- [15] Yaofang Liu, Xiaodong Cun, Xuebo Liu, Xintao Wang, Yong Zhang, Haoxin Chen, Yang Liu, Tieyong Zeng, Raymond Chan, and Ying Shan. 2024. EvalCrafter: Benchmarking and Evaluating Large Video Generation Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 22139–22149. doi:10.1109/CVPR52733.2024.02090
- [16] Yang Liu, Dan Iter, Yichong Xu, et al. 2023. G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment. *CoRR* abs/2303.16634 (2023). <https://arxiv.org/abs/2303.16634>
- [17] Yuanxin Liu, Lei Li, Shuhuai Ren, Rundong Gao, Shicheng Li, Sishuo Chen, Xu Sun, and Lu Hou. 2023. FETV: A Benchmark for Fine-Grained Evaluation of Open-Domain Text-to-Video Generation. In *Advances in Neural Information Processing Systems*.
- [18] Yibo Miao, Yifan Zhu, Yinpeng Dong, Lijia Yu, Jun Zhu, and Xiao-Shan Gao. 2024. T2VSafetyBench: Evaluating the Safety of Text-to-Video Generative Models. *CoRR* abs/2407.05965 (2024). doi:10.48550/arXiv.2407.05965
- [19] Zelu Qi, Ping Shi, Shuqi Wang, et al. 2025. T2VEval: T2V-generated Videos Benchmark Dataset and Objective Evaluation Method. *CoRR* abs/2501.08545 (2025). <https://arxiv.org/abs/2501.08545>
- [20] Qwen Team. 2025. Qwen3-VL Technical Report. *arXiv preprint arXiv:2511.21631* (2025).
- [21] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. 2021. Learning Transferable Visual Models From Natural Language Supervision. In *Proceedings of the 38th International Conference on Machine Learning*. 8748–8763.
- [22] Haoyuan Shi, Yunxin Li, Nanhao Deng, Zhenran Xu, Xinyu Chen, Longyue Wang, Baotian Hu, and Min Zhang. 2026. MSVBench: Towards Human-Level Evaluation of Multi-Shot Video Generation. *arXiv preprint arXiv:2602.23969* (2026).
- [23] Kaiyue Sun, Kaiyi Huang, Xian Liu, Yue Wu, Zihan Xu, Zhenguo Li, and Xihui Liu. 2025. T2V-CompBench: A Comprehensive Benchmark for Compositional Text-to-video Generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 8406–8416. doi:10.1109/CVPR52734.2025.00787
- [24] Tencent. 2025. HunyuanVideo-1.5. <https://huggingface.co/tencent/HunyuanVideo-1.5>.
- [25] Wan Team, Ang Wang, Baole Ai, Bin Wen, Chaojie Mao, Chen-Wei Xie, Di Chen, Feiwu Yu, Haiming Zhao, Jianxiao Yang, Jianyuan Zeng, Jiayu Wang, Jingfeng Zhang, Jingren Zhou, Jinkai Wang, Jixuan Chen, Kai Zhu, Kang Zhao, Keyu Yan, Lianghua Huang, Mengyang Feng, Ningyi Zhang, Pandeng Li, Pingyu Wu, Ruihang Chu, Ruili Feng, Shiwei Zhang, Siyang Sun, Tao Fang, Tianxing Wang, Tianyi Gui, Tingyu Weng, Tong Shen, Wei Lin, Wei Wang, Wenmeng Zhou, Wenten Wang, Wenting Shen, Wenyuan Yu, Xianzhong Shi, Xiaoming Huang, Xin Xu, Yan Kou, Yangyu Lv, Yifei Li, Yijing Liu, Yiming Wang, Yingya Zhang, Yitong Huang, Yong Li, You Wu, Yu Liu, Yulin Pan, Yun Zheng, Yuntao Hong, Yupeng Shi, Yutong Feng, Zeyinzi Jiang, Zhen Han, Zhi-Fan Wu, and Ziyu Liu. 2025. Wan: Open and Advanced Large-Scale Video Generative Models. *CoRR* abs/2503.20314 (2025). doi:10.48550/arXiv.2503.20314
- [26] Weijia Wu, Mingyu Liu, Zeyu Zhu, Xi Xia, Haoen Feng, Wen Wang, Kevin Qinghong Lin, Chunhua Shen, and Mike Zheng Shou. 2025. MovieBench: A Hierarchical Movie Level Dataset for Long Video Generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 28984–28994. doi:10.1109/CVPR52734.2025.02699
- [27] xAI. 2026. Grok Imagine Video 1.5. <https://docs.x.ai/developers/models/grok-imagine-video-1.5>.
- [28] Shenghai Yuan, Jinfa Huang, Yongqi Xu, Yaoyang Liu, Shaofeng Zhang, Yujun Shi, Ruijie Zhu, Xinhua Cheng, Jiebo Luo, and Li Yuan. 2024. ChronoMagic-Bench: A Benchmark for Metamorphic Evaluation of Text-to-Time-lapse Video Generation. In *Advances in Neural Information Processing Systems*.

Table 6: Evaluated systems and the adaptations needed to run the shared StoryEval segment plan.

System	Access	T2V	I2V	Cont. reuse	Adaptation
JoyAI-Echo [6]	Open weights	✓	–	No (T2V only)	Rewrite prompts to the official long-story JSON format
LTX-2.3 [13]	Open weights	✓	✓	Yes	Direct use of canonical fields
HunyuanVideo-1.5 [24]	Open weights	✓	✓	Yes	Direct use of canonical fields
Seedance 2.0 [4]	Closed API	✓	✓	Yes	Alias replacement; anime style for privacy; disable audio when needed
HappyHorse 1.1 [1]	Closed API	✓	✓	Yes	Alias replacement for copyright-sensitive entities
Grok Imagine Video 1.5 [27]	Closed API	–	✓	Yes	T2I start frame via grok-imagine-image-quality; safety softening; DeepSeek-V4-Pro [5] compression if > 4096 characters

A Evaluated Systems and Generation Adaptations

Table 6 lists the six evaluated systems. Open-weight models are run from locally deployed checkpoints; closed models are accessed through official APIs. Continuous boundaries reuse the previous clip’s last frame only when the system supports I2V; otherwise generation falls back to pure T2V.

JoyAI-Echo has no I2V interface, so every segment—including continuous boundaries—is generated by pure T2V after rewriting the prompt into the official long-story template, while preserving the intended events and generation constraints. LTX-2.3 and HunyuanVideo-1.5 consume the canonical fields directly through their native T2V/I2V interfaces. For Seedance 2.0, HappyHorse 1.1, and Grok Imagine Video 1.5, we apply safety and rights mitigations required by the APIs: entity aliases for copyright-sensitive names, anime-style prompting when real-person likeness is rejected, and audio disabled when Seedance flags unsafe soundtrack

content. Because Grok Imagine Video 1.5 is I2V-only, the first segment and segments following narrative transitions first call grok-imagine-image-quality on the planned start-frame description and character visual cues, then animate the resulting image; prompts exceeding the 4096-character limit are compressed by DeepSeek-V4-Pro [5] without changing the intended events.

B Per-System Core Genre×Metric Heatmaps

Figure 3(a) in the main text averages Core scores along the genre or metric axis. Figure 7 reports the full system-wise genre×metric Core heatmaps under a shared color scale, so that every genre–metric cell remains inspectable. Metric abbreviations follow Section 3.3; genre abbreviations follow Section 4. These panels corroborate the main-text claim that SEC/CEC/SV remain comparatively high while TS and KEA remain systematically lower across genres, with HunyuanVideo-1.5 showing a near-uniform low SQS band and closed models remaining densely high overall.




Figure 7: Per-system Core heatmaps over 12 genres (shared scale; Core-item means). Top: HappyHorse, Grok, Seedance. Bottom: Hunyuan, LTX, JoyAI. Abbreviations follow Sections 5.1, 4, and 3.3.

C Additional Adaptive Butterfly Plots

Figure 3(c) in the main text contrasts HappyHorse and Hunyuan. Figure 8 provides the same per-metric Adaptive detail layout for JoyAI-Echo, LTX-2.3, Seedance 2.0, and Grok Imagine Video 1.5. Hard Adaptive blocks remain concentrated on AAC/AFL/OIN, while SCT and related continuity metrics stay comparatively tall when

selected. Beyond the abbreviations defined in Section 5.3, the Adaptive panels also use Intra-Causal Realism (ICR), Character Interaction (CIN), Emotion Dynamics (EMD), Plot-State Fidelity (PSF), Pose Continuity (PCT), Emotion Continuity (ECT), Inter-Causal Retention (ICN), Item Ownership Continuity (IOC), and Identity Continuity (IDC).




Figure 8: Adaptive metric detail for JoyAI, LTX, Seedance, and Grok (same encoding as Figure 3(c)).

D Segment Planning Details

This appendix provides implementation details for the short-drama segment plans used in StoryEval. Each segment plan contains the information required for text-to-video generation, intra-video evaluation, and inter-video continuity evaluation.

D.1 Two-Stage LLM Planning

Given a source excerpt, an LLM first extracts an ordered event trajectory $N = (n_1, \dots, n_T)$. Each node stores a sequential id, a concise event name, and a verbatim continuous source span; the extractor may merge multiple sentences that complete one event but may not rewrite, expand, or reorder the source. Planning then uses two separate LLM calls: the first outputs only the segment count and ordered segment_node_mapping, and the second instantiates the complete plans conditioned on the approved mapping. Thus, trajectory extraction, mapping, and plan instantiation comprise three calls in total. All nodes must be covered in order. A segment normally covers one or two consecutive nodes and may cover three only when they form one continuous action chain that fits the duration constraint. A complex node may recur across consecutive segments only when each occurrence identifies a distinct, continuous sub-event range.

Segments target 5–10 seconds, reflecting the common native duration range of current video generation systems. Abstract narration may be translated into filmable cues, but such optional cinematic choices are generation guidance rather than mandatory source-faithfulness checks. Construction follows the stage-specific criteria in Appendix H.4; flagged outputs are revised or regenerated, but we do not exhaustively audit every boundary or evaluation item.

D.2 Segment Plan Schema

Each short-drama segment plan contains three groups of fields: coverage fields, generation prompt fields, and boundary fields. Coverage fields align the segment with the source-grounded narrative event trajectory. Generation prompt fields specify the visual and narrative content used to generate the video segment. Boundary fields specify how the segment connects to adjacent segments and support inter-segment continuity evaluation. Table 7 summarizes the segment schema.

D.3 Boundary Annotation Rules

StoryEval uses boundary annotations to distinguish frame-level visual continuity from narrative-level transitions. This distinction

prevents reasonable scene changes or time jumps from being incorrectly penalized as continuity failures. Table 8 defines the two boundary types and connects each type to its corresponding evaluation focus. The first segment has no incoming boundary and is excluded from boundary-level evaluation.

Continuous frame anchor checklist. Use this boundary when the same character or key object remains the focus, the previous action naturally continues, the story remains in the same location and time, or preserving the previous end frame helps evaluate character, object, scene, or action continuity. The next segment’s opening 1–2 seconds should preserve the anchored subjects, scene, and object states before the action develops.

Narrative transition checklist. Use this boundary when the next segment starts in a different location or time, shifts to another character viewpoint or information line, cannot naturally reuse the previous end frame, or preserves causal, emotional, or informational continuity rather than visual continuity.

Non-penalization rule. A scene change, time jump, or viewpoint shift should not be treated as a continuity failure under a narrative transition boundary. Conversely, under a continuous frame anchor boundary, unexplained changes in character identity, costume, key object state, scene layout, or action state should be treated as potential continuity errors.

D.4 Model-Specific Sequential Generation

The canonical generation input combines drama_segment with the fields under text_to_video_prompt and replaces the provisional character-consistency text with node-specific visual descriptions from G_s^{gen} . This input is used directly unless a system requires format restructuring or token-budget shortening; such transformations preserve the intended events and generation constraints and are reported per system. When reuse_previous_end_frame is false, a T2V-capable system generates directly from text; an I2V-only system first uses a model-family T2I component to produce the planned start frame. When the flag is true and the system supports I2V, the actual last frame of the previous generated video is passed to the I2V interface; T2V-only systems (e.g., JoyAI-Echo) ignore frame reuse and generate every segment by pure T2V. All systems

Table 7: Segment plan schema used by StoryEval.

Field	Description
Coverage fields	
video_id	Unique identifier of the video segment, e.g., V1 or V2.
covered_nodes	Event node ids covered by the segment. A node id may appear in multiple segments if different subranges are covered.
covered_node_subrange	Textual description of the specific sub-event range covered by the segment.
node_completion_status	Whether the covered node content is partially covered or completed in the current segment.
core_event	One-sentence description of the main event presented in the segment.
segment_function	Narrative function of the segment, such as danger setup, information revelation, emotional turn, action progression, or ending hook.
opening_continuity / ending_hook	Expected opening state and the final narrative or visual state that leads to the next segment.
Generation prompt fields	
suggested_duration_seconds	Suggested video duration, usually between 5 and 10 seconds.
drama_segment	Filmable micro-narrative containing the segment’s visible beginning, action development, and endpoint.
character_visual_description	Character identity and node-specific visible state retrieved from the generation-side story asset view.
scene_description	Location, time, environment, lighting, weather, and period atmosphere.
action_description	Continuous action chain from the first frame to the last frame.
emotional_atmosphere	Visible emotional expression and overall mood of the segment.
camera_language	Concise camera guidance, such as shot scale, angle, and camera movement.
negative_prompt	Constraints used to reduce generation artifacts, such as face changes, extra characters, bad anatomy, modern objects, subtitles, or watermarks.
Boundary fields	
boundary_type	For segments after the first, the relation from the previous segment: continuous frame anchor or narrative transition. The first segment has no incoming boundary.
boundary_reason	Source-grounded reason for assigning the boundary type.
reuse_previous_end_frame	Boolean flag: true for a continuous-frame anchor and false for a narrative transition or the first segment.
continuity_evaluation_targets	Allowed continuity checks, such as character identity, object state, action continuation, narrative causality, information, or emotion.
start_frame	Expected first frame description. For continuous frame anchors, this reuses the previous segment’s end frame.
end_frame	Expected final frame description, which may serve as the next segment’s visual anchor.
continuity_anchor	Structured anchor containing anchor_type, anchor_subjects, anchor_visibility, anchor_state, and reuse_instruction.

Table 8: Boundary annotation types and their evaluation focus.

Boundary type	When to use	Evaluation focus
continuous_frame_anchor	Two adjacent segments continue the same time, space, character, key object, or action chain. The next start frame reuses the previous end frame.	Frame-level continuity of character identity, costume, object state, scene, action, and generation stability.
narrative_transition	The story moves across location, time, character viewpoint, or information line, or strict frame reuse would create an unnatural transition.	Narrative causality, story-state inheritance, information continuation, and emotional direction.

are configured to generate at 720p; model versions and inference settings are recorded in the experimental protocol.

Table 9: Intra-video ontology objects, attributes, and example metrics.

Domain	Object	Attributes	Example metrics
Intra-video	Character	Identity, visual presentation, action/posture, relationship interaction, story state	Character identification; character visual consistency; character action accuracy
Intra-video	Event	Node coverage, node order, event content, causal logic, physical logic	Node completeness; event accuracy; intra-segment causal correctness
Intra-video	Scene	Location, temporal environment, spatial relation, atmosphere	Scene-location consistency; spatial relation plausibility
Intra-video	Key Object	Presence, visual appearance, story state, interaction relation	Object presence; object appearance correctness; object state correctness
Intra-video	Emotion Performance	Emotion type, intensity, change, performance cues	Emotion type matching; emotion change clarity
Intra-video	Cinematic Presentation	Subject visibility, action visibility, object visibility, composition	Subject visibility; action recognizability
Intra-video	Style/Period	Setting consistency, genre style, heterogeneous element control, visual unity	Costume-prop setting consistency; visual style unity
Intra-video	Generation Stability	Character stability, subject integrity, prop stability, frame stability, interference elements	Character generation stability; frame stability

Table 10: Inter-video continuity ontology objects, attributes, and example metrics.

Domain	Object	Attributes	Example metrics
Inter-video	Character Continuity	Identity continuity, pose-position continuity, story-state continuity	Character identity continuity; character state continuity
Inter-video	Object Continuity	Presence continuity, state continuity, spatial relation continuity, possession continuity	Object presence continuity; object state continuity
Inter-video	Scene Continuity	Location continuity, temporal environment continuity, spatial direction continuity	Scene-location continuity; lighting continuity
Inter-video	Action Continuity	Action continuation, action direction, action result, action reset	Action continuation naturalness; action result preservation
Inter-video	Causal Continuity	Causal preservation, result preservation, state rollback control	Inter-segment causal preservation; result preservation
Inter-video	Emotional Continuity	Emotion state continuity, intensity continuity, emotional transition plausibility	Emotion state continuity; emotional transition plausibility
Inter-video	Cinematic Transition	Transition smoothness, action-gaze matching, object matching, composition continuity	Transition smoothness; shot-composition continuity
Inter-video	Style Continuity	Visual texture continuity, period-style continuity, genre-style continuity	Visual texture continuity; period-style continuity
Inter-video	Generation Continuity	Character stability, object stability, image-quality stability	Boundary character stability; boundary object stability

E Metric Grounding Details

E.1 Full Evaluation Ontology

StoryEval defines two evaluation domains: intra-video evaluation and inter-video continuity evaluation. Each domain contains evaluation objects, high-level attributes, and candidate metrics. Table 9 summarizes the intra-video ontology used to evaluate individual segments, while Table 10 summarizes the inter-video ontology used to evaluate continuity across adjacent segments.

E.2 Manually Selected Core Metrics

Core metrics provide a common evaluation basis across all genres and systems. They are selected according to three criteria: coverage

of fundamental story-grounded generation requirements, visual observability, and cross-genre comparability. Table 11 lists the eight segment-level and three boundary-level Core metrics together with the rationale for including each one.

E.3 Evidence Unit Schema and Level Assignment

Each evidence unit is a structured signal used for adaptive metric selection. It is not a final evaluation item or a video label. Table 12 specifies the fields stored for each evidence unit, including its source, ontology mapping, observability, and importance.

Evidence levels $L(e) \in \{0, 1, 2, 3\}$ are assigned according to both observability and importance. The lower panel of Table 12 gives

Table 11: Manually selected Core metrics used across genres (abbreviations defined in Section 3.3).

Abbr.	Core metric	Evaluation object	Rationale
CIA	Character Identity Accuracy	Character	Ensures that the required core characters appear with correct roles.
KEA	Key Event Accuracy	Event	Measures whether the segment presents the central event content.
SEC	Scene & Element Consistency	Scene	Checks whether the segment is grounded in the expected story setting.
KOP	Key Object Presence	Key Object	Measures whether required story-critical objects appear.
EMR	Emotion Match Rate	Emotion Performance	Checks whether visible emotional expression matches the segment context.
SV	Subject Visibility	Cinematic Presentation	Ensures that the main subject can be clearly judged.
CEC	Costume & Environment Consistency	Style/Period	Checks whether visual elements fit the story world and period.
VS	Visual Stability	Generation Stability	Captures basic generation stability issues.
TS	Transition Smoothness	Cinematic Transition	Provides a shared check of whether adjacent clips connect without an unintended abrupt break.
VTC	Visual Texture Continuity	Style Continuity	Checks whether rendering texture and overall visual treatment remain compatible across a boundary.
SQS	Splice Quality Stability	Generation Continuity	Detects abrupt quality degradation or instability at the splice.

Table 12: Evidence unit schema and evidence level assignment for adaptive metric grounding.

Evidence unit fields	
Field	Description
node_id / transition_id	The event node or adjacent segment pair from which the evidence is extracted.
ontology_path	The mapped ontology path, including evaluation domain, object, and attribute.
evidence_summary	A concise summary of the evidence content.
evidence_text	The source text span supporting the evidence unit.
boundary_evidence	The boundary-frame evidence supporting the evidence unit.
visual_observability	Whether the evidence can be observed in video, labeled as high, medium, or low.
importance	The importance of the evidence for the node or transition, labeled as core, supporting, or background.
Evidence level assignment	
Level 0: no matching evidence	The metric is not supported by the node or transition.
Level 1: weak evidence	The evidence is background-level, weakly observable, or only indirectly related to the metric.
Level 2: valid evidence	The evidence is supporting or clearly observable and can support metric selection.
Level 3: strong evidence	The evidence is core to the event or transition and has high or medium observability.

the interpretation of each level and shows how evidence strength is discretized before metric selection.

E.4 Adaptive Metric Selection Algorithm

The adaptive metric selection procedure is evidence-based and genre-level. Let U_c^d denote the units of story c in domain d : event nodes for intra-video metrics and planned adjacent boundaries for inter-video metrics. Table 13 presents the complete domain-specific pipeline from collecting unit-level evidence to returning the genre-adaptive metric set.

E.5 Evaluation Item Schema

Each selected Core or Adaptive metric is instantiated into segment-level or boundary-level evaluation items. The same observable-item structure is used in both domains, with boundary items additionally storing the previous and next segment ids and their expected end/start states. Table 14 defines the common fields and item-level scoring rubric.

E.6 Example of Metric Instantiation

For an intra-video segment from *Wu Song Fights the Tiger*, the adaptive metric *character action accuracy* can be instantiated as shown in Table 15. The table illustrates how an abstract metric is converted into a concrete segment-level evaluation goal, observable checks, and score-specific criteria.

Table 13: Evidence-based genre-adaptive metric selection.

Step	Operation
1	For each metric $m \in M$, iterate over stories $c \in C_g$ under genre g .
2	For each unit $u \in U_c^d$, collect $E(u, m)$, whose evidence object matches m and whose attribute belongs to the related attributes of m .
3	Set $R(u, m) = \max_{e \in E(u, m)} L(e)$. If $E(u, m)$ is empty, set $R(u, m) = 0$.
4	Set $S(c, m) = \max_{u \in U_c^d} R(u, m)$, and set $I(c, m) = 1$ if $S(c, m) \geq 2$; otherwise set $I(c, m) = 0$.
5	Compute the genre-level coverage count $C_g(m) = \sum_c I(c, m)$.
6	If $C_g(m) \geq K_{\text{cov}}$, add m to the domain-specific eligible set $M_{g,d}^{\text{eligible}}$.
7	For each eligible metric, compute $D(c, m) = \frac{1}{ U_c^d } \sum_{u \in U_c^d} R^+(u, m)$, where $R^+(u, m) = R(u, m)$ if $R(u, m) \geq 2$, and 0 otherwise.
8	Compute $\bar{D}_g(m) = \text{average}_c D(c, m)$, rank metrics in $M_{g,d}^{\text{eligible}}$, break ties by the fixed ontology order, and return the top N_d metrics as $M_{g,d}^{\text{adaptive}}$.

Table 14: Evaluation item schema and scoring rubric.

Evaluation item fields	
Field	Description
metric_name	Name of the selected metric.
metric_type	Whether the metric is core or adaptive (both may apply to segment or boundary units).
evaluation_domain	Intra-video (segment) or inter-video (boundary) evaluation.
evaluation_object	Main object evaluated by the metric.
segment_context	Current segment id and event range for intra-video items.
boundary_context	For inter-video items: boundary_id, transition_focus, previous/next segment ids, and expected end/start states.
evaluation_goal	Segment- or boundary-specific goal describing what should be judged.
observable_check_items	Concrete visually observable checks.
scoring_rule	Item-specific guidance for assigning 0/1/2 to each observable check.
Three-level scoring rubric	
Score 0	The required content is missing, severely wrong, or not judgeable.
Score 1	The content is partially presented but incomplete, ambiguous, unstable, or mildly incorrect.
Score 2	The content is clearly, correctly, and stably presented.

Table 15: Example metric instantiation card.

Item	Content
Segment	Wu Song confronts and physically controls the tiger.
Evaluation object	Character.
Evaluation goal	Evaluate whether Wu Song performs the expected action of confronting and controlling the tiger in the current segment.
Observable checks	Wu Song is the active subject; his body movement is directed toward the tiger; the action shows physical contact or control; and the action is visually clear enough to judge.
Score 0	Wu Song does not perform the required action, the tiger interaction is absent, or the main action cannot be judged.
Score 1	Wu Song performs a related action, but the contact, direction, or completion of the action is unclear or only partially visible.
Score 2	Wu Song clearly performs the expected action toward the tiger, with visible direction, physical interaction, and sufficient action clarity.

F Multimodal Story Asset Graph Details

F.1 Story Asset Graph Schema

StoryEval constructs a story-level multimodal asset graph for each source story. The graph contains source-grounded event nodes,

semantic cards, video-derived character anchors, and externally generated object reference images. Table 16 describes the node types stored in the graph together with the relation types that connect events, cards, exemplars, and references.

Table 16: Node and edge types in the multimodal story asset graph.

Node types		
Node type	Description	Main fields
Event Node	Source-grounded narrative event node.	Node id, node name, source text span.
Character Card	Character visual identity and node-level visible states.	Character id, name, visual identity layer, visual state layer, linked node ids.
Object Card	Key-object identity and node-level visible states.	Object id, name, identity layer, state layer, linked node ids.
Scene Card	Scene semantics and spatial context.	Scene id, linked node ids, scene description, period context, spatial features, atmosphere.
Character Exemplar	Video-derived character crop used as a VLM attention anchor.	Exemplar image, linked character, source video/frame.
Object Reference Image	Externally generated object image used as visual prior.	Reference image, linked object, generation prompt.
Edge types		
Edge type	Meaning	
character_appears_in_node	A character appears or acts in an event node.	
character_has_visual_state_in_node	A character has a node-specific visible state.	
object_appears_in_node	A key object appears in an event node.	
object_has_visual_state_in_node	A key object has a node-specific visible state.	
scene_covers_node	A scene covers one or more event nodes.	
character_linked_to_exemplar	A character is linked to video-derived exemplar images.	
object_linked_to_reference	A key object is linked to externally generated object reference images.	

F.2 Permission-Controlled Graph Views

The same story asset graph supports generation and evaluation through different views. The generation view G_s^{gen} exposes only character identity (name and stable identity text) and node-specific visible-state descriptions needed to construct `character_visual_description`. It withholds object and scene cards, card reference images, metrics, evaluation items, rubrics, scores, video-derived exemplars, external object references, and evaluator rationales. The evaluation view G_s^{eval} retrieves character, object, and scene cards for the current nodes—including identity/state reference images when available—and may additionally include verified character exemplars and externally generated object references. This separation supports consistent character prompting without treating the graph as a full T2V conditioning signal or leaking evaluation answers.

F.3 Character Exemplar Extraction

Character exemplars are used to help the VLM focus on the correct character during evaluation. They are not used as ground-truth visual targets for judging whether a generated character appearance is correct. The extraction procedure is summarized together with the object reference image generation procedure in Table 17. It shows how sampled frames are converted into verified character exemplars that can be used as evaluation-time attention anchors. These exemplars are used only as evaluation-time attention anchors.

They help the VLM associate actions, interactions, and spatial relations with the correct character in the generated video.

F.4 Object Reference Image Generation

Object reference images are generated externally because key objects are themselves evaluation targets. Directly extracting object images from the evaluated video could weaken the ability to detect missing, incorrect, or anachronistic objects. The generation procedure is summarized together with character exemplar extraction in Table 17. It shows how object-card descriptions are converted into visual priors for VLM scoring.

Object reference images are not used for pixel-level matching. They serve as visual priors for judging whether the generated video contains an object of the correct category, visible form, and historical style.

F.5 Reference Card Schema

Reference cards store the semantic and visual reference information used by the graph. Table 18 summarizes the layers and fields used for character, object, and scene cards. Reference images attached to character/object cards are retrieved only under G_s^{eval} . Generation uses the textual identity and state fields exposed by G_s^{gen} .

F.6 Example Reference Cards

Table 19 provides example reference cards from *Wu Song Fights the Tiger*. The example shows how the same story segment is represented through character, object, and scene references.

Table 17: Reference asset construction pipelines.

Reference asset	Step	Operation
Character exemplar	1	Uniformly sample 5 frames per second from each generated video.
Character exemplar	2	Detect person regions in sampled frames using YOLO26.
Character exemplar	3	Keep the top-10 person candidates according to detection confidence.
Character exemplar	4	Provide the candidate crops and story-level character cards to a VLM.
Character exemplar	5	Ask the VLM to identify which candidate crops correspond to which character.
Character exemplar	6	For each character, retain 1–3 verified exemplar images.
Character exemplar	7	If no reliable crop is found for a character, no exemplar image is assigned.
Object reference	1	Use the object card identity layer and state layer to construct a text-to-image prompt.
Object reference	2	Generate object reference images with an external text-to-image model.
Object reference	3	Link generated images to the corresponding object card.
Object reference	4	Use these images as category, form, and period-style references during VLM scoring.

Table 18: Reference card schema.

Card type	Layer	Fields
Character Card	Visual Identity Layer	Character id, name, stable visual features, period context, prompt; identity reference images (evaluation view only).
Character Card	Visual State Layer	State id, linked node ids, visual change description, edit prompt; state reference images (evaluation view only).
Object Card	Identity Layer	Object id, name, stable visual features, period context, prompt; identity reference images (evaluation view only).
Object Card	State Layer	State id, linked node ids, visual change description, edit prompt; state reference images (evaluation view only).
Scene Card	Semantic Layer	Scene id, linked node ids, scene name, scene description, period context, spatial features, atmosphere.

Table 19: Example reference cards from *Wu Song Fights the Tiger*.

Card	Field	Content
Character	Character	Wu Song.
Character	Visual identity	Strong adult male traveler and martial hero in a classical Chinese setting, wearing ancient traveler clothing.
Character	Linked nodes	Nodes where Wu Song drinks, walks to the ridge, encounters the tiger, and fights it.
Character	Visual state	If no major visible change occurs, the same visual identity is shared across nodes.
Character	Character exemplar	1–3 video-derived crops verified as likely corresponding to Wu Song, used as VLM attention anchors.
Object	Object	Staff.
Object	Visual identity	Long wooden staff with simple ancient travel-weapon style.
Object	Linked nodes	Nodes where Wu Song carries, uses, or loses the staff.
Object	Visual state	A separate state is created only if the staff becomes visibly broken, lost, or otherwise visually changed.
Object	Object reference	Externally generated images of an ancient wooden staff, used as category and period-style references.
Scene	Scene	Jingyang Ridge.
Scene	Scene description	A wild mountain ridge with sparse trees, uneven ground, and an isolated atmosphere.
Scene	Linked nodes	Nodes where Wu Song walks alone, encounters the tiger, and fights it.
Scene	Spatial features	Open mountain path, vegetation as partial occlusion, uneven terrain suitable for physical struggle.
Scene	Atmosphere	Dangerous, tense, and isolated.

G Reference-Guided VLM Scoring Details

Scoring input. For an intra-video item, the VLM receives the generated segment, its metric name and evaluation goal, the ordered

observable checks, the item-specific 0/1/2 guidance, and references retrieved from G_s^{eval} . For an inter-video item, the visual input is the last two seconds of the previous segment and the first two

Table 20: Source works used in StoryEval.

Work	Tradition	Main Coverage	Text Source
Romance of the Three Kingdoms	Chinese classic	Stratagem, Action, Patriotism	Public-domain / verified edition
Water Margin	Chinese classic	Action, Drama, Suspense	Public-domain / verified edition
Journey to the West	Chinese classic	Fantasy, Adventure, Supernatural Tales	Public-domain / verified edition
Dream of the Red Chamber	Chinese classic	Drama, Historical Romance, Female Growth, Palace Intrigue	Public-domain / verified edition
The Odyssey	World classic	Adventure, Fantasy, Action	Project Gutenberg
The Three Musketeers	World classic	Stratagem, Action, Adventure	Project Gutenberg
Pride and Prejudice	World classic	Historical Romance, Comedy, Female Growth	Project Gutenberg
Grimms' Fairy Tales	World classic	Fantasy, Adventure, Comedy, Suspense, Supernatural Tales	Project Gutenberg

seconds of the next segment. Continuous-frame anchors use visual and state continuity checks, whereas narrative transitions exclude strict frame, location, and appearance matching when a change is intended.

Output constraints and aggregation. The judge must return one JSON record per observable check with the unchanged observation text, a score in $\{0, 1, 2\}$, and video-grounded evidence. The reported sum must equal the item-score sum. Unit- and story-level scores follow Equations 1 and 2. Boundary, genre, and system scores average only their corresponding normalized units; system-level tables use equal-weight story macro-averaging rather than pooling all items globally. The scoring prompt instructs the VLM to use graph references as entity and context aids rather than pixel-level targets and to identify the specific mismatch when assigning a low score.

H Benchmark Instantiation Details

This appendix provides additional details for the StoryEval instantiation described in Section 4. We include the source works, the operational definitions of the 12 short-drama genres, the complete list of selected source-story excerpts, and the criteria used when reviewing excerpts and downstream LLM outputs. These details support reproducibility and clarify how the selected excerpts cover both Chinese and world classical narratives.

H.1 Source Works

StoryEval is instantiated from eight classical narrative works, including four Chinese classical novels and four world literary classics. Table 20 summarizes the source works, their narrative traditions, major genre coverage, and text sources. The Chinese works provide dense event structures commonly used in Chinese short-drama genres, such as political strategy, mythological fantasy, household conflict, martial action, and emotional drama. The world classics provide cross-culturally recognizable scenarios, including Adventure, Suspense, Historical Romance, Supernatural Tales, and Stratagem.

H.2 Genre Taxonomy and Functional Mapping

Table 21 provides the operational definitions of the 12 genre labels used in StoryEval, together with the abbreviations introduced in Section 4. These labels are designed to cover common short-drama narrative demands while remaining interpretable across Chinese and world classical sources. For culturally specific categories, we use functional mappings. For example, martial-action scenes are mapped to Action, period love stories to Historical Romance, and uncanny events to Supernatural Tales.

Table 21: Genre taxonomy, abbreviations, and operational definitions.

Functional Category	English Label	Abbr.	Operational Definition
Plot conflict	Drama	Dram	Character-centered conflict, emotional turn, or consequential narrative event.
Strategy and power	Stratagem	Strat	Strategy, deception, political calculation, or manipulation of information.
Physical conflict	Action	Act	Visually salient physical conflict, chase, combat, or forceful movement.
Magical events	Fantasy	Fant	Magical beings, transformation, divine power, or impossible events.
Journey and danger	Adventure	Adv	Journey, exploration, danger, escape, or traversal through unfamiliar spaces.
Comic reversal	Comedy	Com	Misunderstanding, social awkwardness, comic reversal, or exaggerated behavior.
Hidden information	Suspense	Susp	Hidden information, clue discovery, threat uncertainty, or delayed revelation.
Period relationship	Historical Romance	H-Rom	Romantic or relational conflict in a historical/social-period setting.
Female agency	Female Growth	FGrow	Female agency, resistance, self-recognition, or autonomous decision-making.
Household power	Palace Intrigue	Pal-Int	Power struggle within court, family, household, or elite social hierarchy.
Collective duty	Patriotism	Patr	Loyalty, sacrifice, public duty, homeland restoration, or collective cause.
Uncanny event	Supernatural Tales	Sp-Tales	Ghostly, monstrous, uncanny, or supernatural events.

H.3 Selected Source-Story Excerpts

Tables 22–26 list all 60 source-story excerpts used in StoryEval. Each genre contains five excerpts jointly selected from Chinese and world classics. For each excerpt, we report the source work, the selected story event, source tradition, and the main visual elements expected

to support video generation and evaluation. The selected excerpts are intended to provide compact narrative units with clear event structure and strong visual expressiveness. Genre names follow Table 21; for table compactness, longer labels may be shortened as Susp, H-Rom, Pal-Int, Patr, and Sp-Tales.

Table 22: Selected source-story excerpts for Drama, Stratagem, Action, and Fantasy.

Genre	Work	Selected Excerpt	Tradition	Main Visual Elements
Drama	Dream of the Red Chamber	Baoyu is beaten and Daiyu visits him afterward	Chinese	injured Baoyu, Daiyu, bedside scene, tears, restrained emotion
Drama	Water Margin	Lin Chong at the snowstorm mountain temple	Chinese	snowy temple, spear, fire, night, betrayal aftermath
Drama	Dream of the Red Chamber	Daiyu buries fallen flowers	Chinese	Daiyu, flower hoe, fallen petals, flower mound, restrained grief
Drama	Pride and Prejudice	Elizabeth reads Darcy’s letter and reconsiders him	World	letter, Elizabeth alone, changing expression, quiet room
Drama	Grimms’ Fairy Tales	Ashputtel sorts peas from ashes but is still denied the festival	World	ash heap, peas, birds, stepfamily, visible disappointment
Stratagem	Romance of the Three Kingdoms	Borrowing arrows with straw boats	Chinese	fog, straw boats, arrows, river, military deception
Stratagem	Romance of the Three Kingdoms	Jiang Gan steals the forged letter in Zhou Yu’s camp	Chinese	candlelit tent, hidden letter, feigned drunkenness, spying
Stratagem	Romance of the Three Kingdoms	Zhuge Liang’s empty city strategy	Chinese	city gate, zither, suspicious army, calm strategist
Stratagem	The Three Musketeers	Cardinal Richelieu plots around the Queen’s diamond studs	World	court chamber, letter, jewels, Cardinal, political trap
Stratagem	The Three Musketeers	Milady steals Buckingham’s diamond studs	World	jewels, disguise, secret exchange, Milady, political deception
Action	Water Margin	Wu Song fights the tiger	Chinese	tiger, ridge, broken staff, close combat, forest path
Action	Water Margin	The heroes launch the third attack on Zhu Family Manor	Chinese	fortified manor, fighters, weapons, coordinated assault, battlefield movement
Action	Journey to the West	Sun Wukong causes havoc in Heaven	Chinese	heavenly palace, Monkey King, golden cudgel, divine soldiers
Action	The Odyssey	Odysseus blinds the Cyclops	World	cave, giant, wooden stake, fire, escaping sailors
Action	The Three Musketeers	d’Artagnan and the three musketeers fight the guards	World	swords, street duel, four fighters, guards, rapid movement
Fantasy	Journey to the West	The contest of magic in Chechi Kingdom	Chinese	altar, monks, ritual, magic contest, supernatural transformation
Fantasy	Journey to the West	Borrowing the Plantain Fan at the Flaming Mountains	Chinese	magic fan, flames, mountain road, transformation, combat
Fantasy	Journey to the West	Sun Wukong learns supernatural arts	Chinese	master and disciple, mountain cave, magical training, transformation, cloud movement
Fantasy	The Odyssey	Circe transforms sailors into pigs	World	sorceress, cup, palace, pigs, magical transformation
Fantasy	Grimms’ Fairy Tales	The Frog Prince transforms into a prince	World	frog, princess, palace chamber, magical transformation, royal clothing

Table 23: Selected source-story excerpts for Adventure and Comedy.

Genre	Work	Selected Excerpt	Tradition	Main Visual Elements
Adventure	Journey to the West	The pilgrims are blocked by the Flaming Mountains	Chinese	burning mountain, desert road, exhausted travelers, blocked path
Adventure	Journey to the West	Danger at the River of Heaven	Chinese	wide river, boat, water monster threat, pilgrims, stormy water
Adventure	Journey to the West	The pilgrims confront Sha Wujing at the Flowing Sand River	Chinese	dangerous river, sand monk, staff combat, pilgrims, river crossing
Adventure	The Odyssey	Odysseus escapes from the Cyclops' cave	World	cave entrance, sheep, hidden sailors, blinded giant, escape
Adventure	Grimms' Fairy Tales	Hansel and Gretel discover the gingerbread house	World	forest, lost children, edible house, cautious approach, unfamiliar refuge
Comedy	Dream of the Red Chamber	Granny Liu visits the Grand View Garden	Chinese	elderly visitor, noble garden, banquet, laughter, social contrast
Comedy	Journey to the West	Sun Wukong and his companions steal the ginseng fruit	Chinese	magical fruit tree, furtive picking, surprised companions, comic concealment
Comedy	Journey to the West	Zhu Bajie carries the disguised bride	Chinese	pig-like hero, disguised bride, mountain path, exaggerated effort, comic reveal
Comedy	Pride and Prejudice	Mr. Collins proposes to Elizabeth and is rejected	World	formal sitting room, awkward proposal, Elizabeth's refusal, comic stiffness
Comedy	Grimms' Fairy Tales	The Bremen Town Musicians frighten away the robbers	World	animal group, cottage window, startled robbers, chaotic comic attack

Table 24: Selected source-story excerpts for Suspense and Historical Romance.

Genre	Work	Selected Excerpt	Tradition	Main Visual Elements
Suspense	Water Margin	Lin Chong is lured into White Tiger Hall	Chinese	military hall, forbidden weapon, officials, sudden accusation
Suspense	Dream of the Red Chamber	You Sanjie kills herself after Liu Xianglian's suspicion	Chinese	returned sword, confrontation, sudden self-harm, shocked witnesses
Suspense	Grimms' Fairy Tales	Snow White accepts the poisoned apple	World	disguised queen, apple, cottage, hesitation, concealed threat
Suspense	Grimms' Fairy Tales	Hansel and Gretel discover the witch's true intent	World	cage, oven, witch, frightened children, hidden danger
Suspense	Grimms' Fairy Tales	The maiden discovers the Robber Bridegroom's secret	World	hidden witness, robbers, cellar, severed finger clue, escape
Historical Romance	Dream of the Red Chamber	Baoyu and Daiyu meet for the first time	Chinese	aristocratic interior, young nobles, first gaze, jade pendant
Historical Romance	Dream of the Red Chamber	Baoyu gives handkerchiefs and Daiyu writes on them	Chinese	handkerchiefs, brush writing, private emotion, chamber scene
Historical Romance	Dream of the Red Chamber	Daiyu visits Baoyu after he is beaten	Chinese	bedside, restrained affection, tears, injury, quiet room
Historical Romance	Pride and Prejudice	Darcy's first proposal to Elizabeth is rejected	World	proposal, tense conversation, drawing room, emotional reversal
Historical Romance	Pride and Prejudice	Elizabeth unexpectedly meets Darcy again at Pemberley	World	estate grounds, unexpected encounter, restrained expressions, period clothing

Table 25: Selected source-story excerpts for Female Growth and Palace Intrigue.

Genre	Work	Selected Excerpt	Tradition	Main Visual Elements
Female Growth	Dream of the Red Chamber	Tanchun manages the household affairs	Chinese	account books, servants, authoritative young woman, household order
Female Growth	Dream of the Red Chamber	Yuanyang resists forced marriage	Chinese	family pressure, refusal, confrontation, servant woman's agency
Female Growth	Dream of the Red Chamber	Baochai prepares to leave her boudoir	Chinese	young woman, packed belongings, chamber, farewell, independent decision
Female Growth	Pride and Prejudice	Elizabeth rejects Mr. Collins's proposal	World	formal proposal, firm refusal, social pressure, self-possession
Female Growth	Pride and Prejudice	Elizabeth resists Lady Catherine's pressure	World	confrontation, aristocratic pressure, Elizabeth standing firm
Palace Intrigue	Dream of the Red Chamber	Wang Xifeng manipulates affairs at Iron Threshold Temple	Chinese	temple, documents, secret arrangement, Wang Xifeng, bribery implication
Palace Intrigue	Dream of the Red Chamber	Wang Xifeng sets the trap of unrequited love	Chinese	private room, manipulated lovers, letter/message, hidden scheme
Palace Intrigue	Dream of the Red Chamber	The search of Grand View Garden	Chinese	household inspection, maids, authority figures, hidden belongings
Palace Intrigue	Romance of the Three Kingdoms	Diaochan's linked stratagem against Dong Zhuo and Lü Bu	Chinese	court chamber, Diaochan, two powerful men, staged intimacy
Palace Intrigue	The Three Musketeers	Milady poisons Constance	World	guarded room, poison, Milady, Constance, concealed assassination

Table 26: Selected source-story excerpts for Patriotism and Supernatural Tales.

Genre	Work	Selected Excerpt	Tradition	Main Visual Elements
Patriotism	Romance of the Three Kingdoms	The oath of the Peach Garden	Chinese	three sworn brothers, peach garden, ritual, weapons, loyalty
Patriotism	Romance of the Three Kingdoms	Zhao Yun rescues Liu Bei's infant at Changban	Chinese	battlefield, infant, horse, spear, pursuing soldiers
Patriotism	Romance of the Three Kingdoms	Huang Gai's fire ships strike Cao Cao's fleet at Red Cliffs	Chinese	fire ships, river fleet, soldiers, flames, collective battle
Patriotism	The Odyssey	Odysseus returns home and restores order	World	hall, bow, suitors, household restoration, returning king
Patriotism	The Three Musketeers	The musketeers undertake the mission to protect the Queen's honor	World	royal token, journey preparation, swords, loyal companions
Supernatural Tales	Journey to the West	Sun Wukong exposes the White Bone Demon	Chinese	demon disguise, skeletal reveal, Monkey King, eerie landscape
Supernatural Tales	Journey to the West	The true and false Monkey Kings confront each other	Chinese	two identical Monkey Kings, staff duel, divine judgment, identity confusion
Supernatural Tales	Journey to the West	The Spider Demons trap the pilgrims in the cave	Chinese	spider webs, cave, demon women, trapped monk, uncanny threat
Supernatural Tales	Grimms' Fairy Tales	Snow White collapses after eating the poisoned apple	World	apple, cottage, Snow White, sudden collapse, disguised queen
Supernatural Tales	Grimms' Fairy Tales	Hansel and Gretel push the witch into the oven	World	oven, witch, children, fire, supernatural threat defeated

H.4 Instantiation Quality Control

Table 27 lists the criteria used to guide revision of excerpts, trajectories, plans, evaluation items, and story assets. These criteria are not a claim that every field was independently human-verified; as noted in Section 4, we do not exhaustively audit every boundary or evaluation item.

Table 27: Quality control criteria for StoryEval instantiation.

Criterion	Check
Genre fit	The excerpt clearly exhibits the target genre features.
Event closure	The excerpt contains a compact local narrative unit.
Visual salience	The excerpt includes observable characters, objects, actions, or scenes.
Segmentability	The excerpt can be decomposed into short video segments.
Source fidelity	The selected event remains grounded in the original story.
Cultural clarity	The excerpt is understandable with minimal background knowledge.
Evaluation value	The excerpt supports observable diagnosis of event realization or cross-segment continuity.
Event trajectory	Nodes use verbatim source spans, preserve narrative order, and cover the selected excerpt without unsupported additions.
Segment mapping	Every node is covered in order; repeated nodes use distinct continuous subranges; each segment is feasible within 5–10 seconds.
Boundary annotation	Narrative-transition boundaries set frame reuse to false; continuous-frame anchors set it to true and have a valid visible anchor. The first segment has no incoming boundary.
Evaluation items	Metrics remain unchanged, checks are visually observable, and optional cinematic details are not treated as mandatory source evidence.
Story assets	Character, object, and scene cards match their linked nodes; generation and evaluation views expose only their permitted fields.